


Red Black Trees



Today

- Balanced search trees, or how to avoid this even in the worst case





Red-black trees


BSTs with an extra one-bit **color** field in each node.

Red-black properties:

1. Every node is either red or black.
2. The root and leaves (**NIL**'s) are black.
3. If a node is red, then its parent is black.
4. All simple paths from any node **x** to a descendant leaf have the same number of black nodes.



Example of a red-black tree






Use of red-black trees

- What properties would we like to prove about red-black trees ?
 - They **always** have $O(\log n)$ height
 - There is an $O(\log n)$ –time insertion procedure which preserves the red-black properties
- Is it true that, after we add a new element to a tree (as in the previous lecture), we can always recolor the tree to keep it red-black ?




Example of a red-black tree






Use of red-black trees

- What properties would we like to prove about red-black trees ?
 - They **always** have $O(\log n)$ height
 - There is an $O(\log n)$ –time insertion procedure which preserves the red-black properties
- Is it true that, after we add a new element to a tree (as in the previous lecture), we can always recolor the tree to keep it red-black ?
- NO
- After insertions, sometimes we need to juggle nodes around




Rotations






Rotations maintain the inorder ordering of keys:

- $a \in \alpha, b \in \beta, c \in \gamma \Rightarrow a \leq A \leq b \leq B \leq c.$

A rotation can be performed in $O(1)$ time.



Rotations can reduce height





Red-black tree wrap-up

- Can show how
 - $O(\log n)$ re-colorings
 - 1 rotationcan restore red-black properties after an insertion
- Instead, we will see 2-3 trees (but will come back to red-black trees at the end)




Height of a red-black tree

Theorem. A red-black tree with n keys has height $h \leq 2 \lg(n + 1)$.

INTUITION:

- Merge red nodes into their black parents.






Height of a red-black tree

Theorem. A red-black tree with n keys has height $h \leq 2 \lg(n + 1)$.

INTUITION:

- Merge red nodes into their black parents.






Height of a red-black tree

Theorem. A red-black tree with n keys has height
 $h \leq 2 \lg(n + 1)$.

INTUITION:

- Merge red nodes into their black parents.






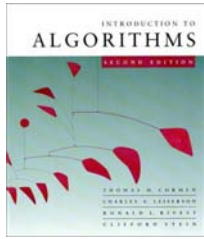
Height of a red-black tree

Theorem. A red-black tree with n keys has height
 $h \leq 2 \lg(n + 1)$.

INTUITION:

- Merge red nodes into their black parents.







Height of a red-black tree

Theorem. A red-black tree with n keys has height
 $h \leq 2 \lg(n + 1)$.

INTUITION:

- Merge red nodes into their black parents.





Red-Black Trees

Red-black trees

- A variation of binary search trees.
- **Balanced**: height is $O(\lg n)$, where n is the number of nodes.
- Operations will take $O(\lg n)$ time in the worst case.




Red-Black Trees

A *red-black tree* is a binary search tree + 1 bit per node: an attribute *color*, which is either **red** or **black**.

All leaves are empty (nil) and colored black.

- We use a single sentinel, $nil[T]$, for all the leaves of red-black tree T .
- $color[nil[T]]$ is black.
- The root's parent is also $nil[T]$.


All other attributes of binary search trees are inherited by red-black trees (*key*, *left*, *right*, and *p*). We don't care about the key in $nil[T]$.



Red-Black Trees


Red-black properties

1. Every node is either **red** or **black**.
2. The root is black.
3. Every leaf ($nil[T]$) is black.
4. If a node is red, then both its children are black. (Hence no two reds in a row on a simple path from the root to a leaf.)
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes.




Red-Black Trees

Example:



[Nodes with bold outline indicate black nodes. Don't add heights and black-heights yet. We won't bother with drawing $nil[T]$ any more.]



Red-Black Trees

Height of a red-black tree

- *Height of a node* is the number of edges in a longest path to a leaf.
- *Black-height* of a node x : $bh(x)$ is the number of black nodes (including $nil[T]$) on the path from x to leaf, not counting x . By property 5, black-height is well defined.



Red-Black Trees

Claim

Any node with height h has black-height $\geq h/2$.

Proof By property 4, $\leq h/2$ nodes on the path from the node to a leaf are red.
Hence $\geq h/2$ are black. (claim)

Claim


The subtree rooted at any node x contains $\geq 2^{\text{bh}(x)} - 1$ internal nodes.

Proof By induction on height of x .

Basis: Height of $x = 0 \Rightarrow x$ is a leaf $\Rightarrow \text{bh}(x) = 0$.

The subtree rooted at x has 0 internal nodes. $2^0 - 1 = 0$.

Inductive step: Let the height of x be h and $\text{bh}(x) = b$. Any child of x has height $h - 1$ and black-height either b (if the child is red) or $b - 1$ (if the child is black). By the inductive hypothesis, each child has $\geq 2^{\text{bh}(x)-1} - 1$ internal nodes. Thus, the subtree rooted at x contains $\geq 2 \cdot (2^{\text{bh}(x)-1} - 1) + 1 = 2^{\text{bh}(x)} - 1$ internal nodes. (The +1 is for x itself.) (claim)



Red-Black Trees


Lemma

A red-black tree with n internal nodes has height $\leq 2 \lg(n + 1)$.

Proof Let h and b be the height and black-height of the root, respectively. By the above two claims,

$$n \geq 2^b - 1 \geq 2^{h/2} - 1 .$$

Adding 1 to both sides and then taking logs gives $\lg(n + 1) \geq h/2$, which implies that $h \leq 2 \lg(n + 1)$. (theorem)



Red-Black Trees

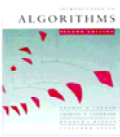
Operations on red-black trees

The non-modifying binary-search-tree operations MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR, and SEARCH run in $O(\text{height})$ time. Thus, they take $O(\lg n)$ time on red-black trees.

Insertion and deletion are not so easy.

If we **insert**, what color to make the new node?

- **Red?** Might violate property 4.
- **Black?** Might violate property 5.



Red-Black Trees

If we **delete**, thus removing a node, what color was the node that was removed?

- **Red?** OK, since we won't have changed any black-heights, nor will we have created two red nodes in a row. Also, cannot cause a **violation of property 2**, since if the removed node was red, it could not have been the root.
- **Black?** Could cause there to be two reds in a row (**violating property 4**), and can also cause a **violation of property 5**. Could also cause a **violation of property 2**, if the removed node was the root and its child—which becomes the new root—was red.




Rotations

Rotations

- The basic tree-restructuring operation.
- Needed to maintain red-black trees as **balanced binary search trees**.
- Changes the local pointer structure. (Only pointers are changed.)
- Won't upset the binary-search-tree property.
- Have both **left rotation** and **right rotation**.
They are inverses of each other.
- A rotation takes a red-black-tree and a node within the tree.

Rotations





LEFT-ROTATE(T, x)

.....>>>

<<<.....

RIGHT-ROTATE(T, y)





Rotations

LEFT-ROTATE(T, x)

$y \leftarrow \text{right}[x]$ \triangleright Set y .

$\text{right}[x] \leftarrow \text{left}[y]$ \triangleright Turn y 's left subtree into x 's right subtree.

if $\text{left}[y] \neq \text{nil}[T]$

then $p[\text{left}[y]] \leftarrow x$

$p[y] \leftarrow p[x]$ \triangleright Link x 's parent to y .

if $p[x] = \text{nil}[T]$

then $\text{root}[T] \leftarrow y$

else if $x = \text{left}[p[x]]$

then $\text{left}[p[x]] \leftarrow y$

else $\text{right}[p[x]] \leftarrow y$

$\text{left}[y] \leftarrow x$ \triangleright Put x on y 's left.


$p[x] \leftarrow y$

Pseudocode for RIGHT-ROTATE is symmetric: exchange *left* and *right* everywhere.




Rotations

Example: [Use to demonstrate that rotation maintains inorder ordering of keys. Node colors omitted.]



Time: $O(1)$ for both LEFT-ROTATE and RIGHT-ROTATE, since a constant number of pointers are modified.



Insertion

Insertion Start by doing regular binary-search-tree insertion:

RB-INSERT(T, z)

$y \leftarrow \text{nil}[T]$

$x \leftarrow \text{root}[T]$

while $x \neq \text{nil}[T]$

do $y \leftarrow x$

if $\text{key}[z] < \text{key}[x]$

then $x \leftarrow \text{left}[x]$

else $x \leftarrow \text{right}[x]$

$p[z] \leftarrow y$

if $y = \text{nil}[T]$

then $\text{root}[T] \leftarrow z$

else if $\text{key}[z] < \text{key}[y]$

then $\text{left}[y] \leftarrow z$


else $\text{right}[y] \leftarrow z$

$\text{left}[z] \leftarrow \text{nil}[T]$

$\text{right}[z] \leftarrow \text{nil}[T]$

$\text{color}[z] \leftarrow \text{RED}$

RB-INSERT-FIXUP(T, z)




Insertion

- RB-INSERT ends by coloring the new node z red.
- Then it calls **RB-INSERT-FIXUP** because we could have violated a red-black property.

Which property might be violated?

1. OK.
2. If z is the root, then there's a violation. Otherwise, OK.
3. OK.
4. If $p[z]$ is red, there's a violation: both z and $p[z]$ are red.
5. OK.

Remove the violation by calling **RB-INSERT-FIXUP**:



Insertion

RB-INSERT-FIXUP(T, z)

while $color[p[z]] = \text{RED}$

do if $p[z] = \text{left}[p[p[z]]]$

then $y \leftarrow \text{right}[p[p[z]]]$

if $color[y] = \text{RED}$

then $color[p[z]] \leftarrow \text{BLACK}$

▷ Case 1

$color[y] \leftarrow \text{BLACK}$

▷ Case 1

$color[p[p[z]]] \leftarrow \text{RED}$

▷ Case 1

$z \leftarrow p[p[z]]$

▷ Case 1

else if $z = \text{right}[p[z]]$

then $z \leftarrow p[z]$

▷ Case 2

 LEFT-ROTATE(T, z)

▷ Case 2

$color[p[z]] \leftarrow \text{BLACK}$

▷ Case 3

$color[p[p[z]]] \leftarrow \text{RED}$

▷ Case 3

 RIGHT-ROTATE($T, p[p[z]]$)

▷ Case 3

else (same as **then** clause





 with “right” and “left” exchanged)

$color[\text{root}[T]] \leftarrow \text{BLACK}$



Insertion

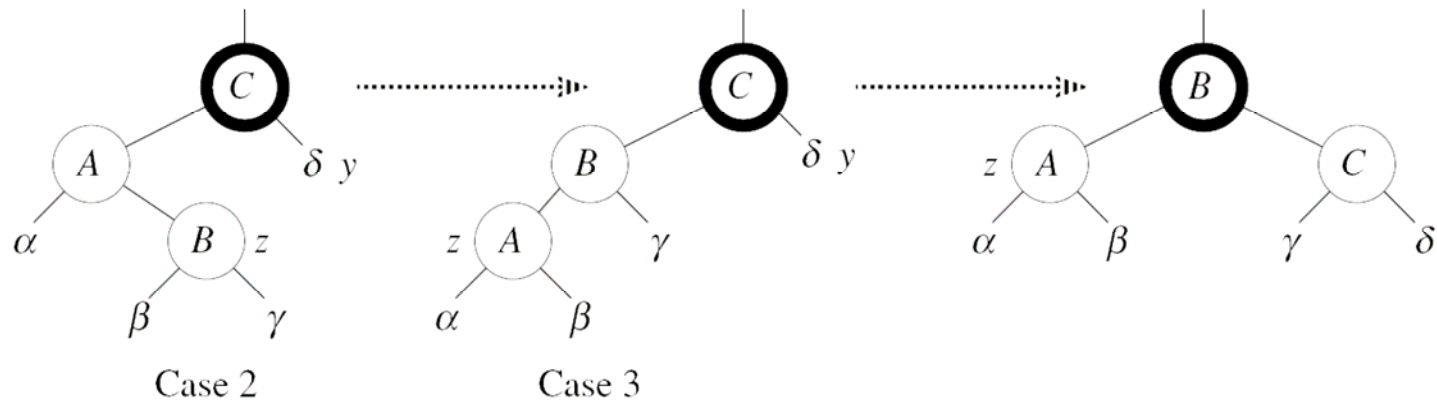
Case 1: y is red



- $p[p[z]]$ (z 's grandparent) must be black, since z and $p[z]$ are both red and there are no other violations of property 4.
- Make $p[z]$ and y black \Rightarrow now z and $p[z]$ are not both red. But property 5 might now be violated.
- Make $p[p[z]]$ red \Rightarrow restores property 5.
- The next iteration has $p[p[z]]$ as the new z (i.e., z moves up 2 levels).

Insertion


Case 2: y is black, z is a right child



- Left rotate around $p[z] \Rightarrow$ now z is a left child, and both z and $p[z]$ are red.
- Takes us immediately to case 3.

Case 3: y is black, z is a left child

- Make $p[z]$ black and $p[p[z]]$ red.
- Then right rotate on $p[p[z]]$.
- No longer have 2 reds in a row.
- $p[z]$ is now black \Rightarrow no more iterations.



Insertion

Analysis

$O(\lg n)$ time to get through RB-INSERT up to the call of RB-INSERT-FIXUP.

Within RB-INSERT-FIXUP:

- Each iteration takes $O(1)$ time.
- Each iteration is either the last one or it moves z up 2 levels.
- $O(\lg n)$ levels $\Rightarrow O(\lg n)$ time.
- Also note that there are at most 2 rotations overall.

Thus, **insertion** into a red-black tree takes $O(\lg n)$ time.




Deletion

Start by doing regular binary-search-tree deletion:

```
RB-DELETE( $T, z$ )
if  $left[z] = nil[T]$  or  $right[z] = nil[T]$ 
    then  $y \leftarrow z$ 
    else  $y \leftarrow \text{TREE-SUCCESSOR}(z)$ 
if  $left[y] \neq nil[T]$ 
    then  $x \leftarrow left[y]$ 
    else  $x \leftarrow right[y]$ 
 $p[x] \leftarrow p[y]$ 
if  $p[y] = nil[T]$ 
    then  $root[T] \leftarrow x$ 
    else if  $y = left[p[y]]$ 
        then  $left[p[y]] \leftarrow x$ 
        else  $right[p[y]] \leftarrow x$ 
if  $y \neq z$ 
    then  $key[z] \leftarrow key[y]$ 
        copy  $y$ 's satellite data into  $z$ 
if  $color[y] = \text{BLACK}$ 
    then RB-DELETE-FIXUP( $T, x$ )
return  $y$ 
```

- y is the node that was actually spliced out.
 - x is either
 - y 's sole non-sentinel child before y was spliced out, or
 - the sentinel, if y had no children.
- In both cases, $p[x]$ is now the node that was previously y 's parent.



Deletion

If y is black, we could have violations of red-black properties:

1. OK.
2. If y is the root and x is red, then the root has become red.
3. OK.
4. Violation if $p[y]$ and x are both red.
5. Any path containing y now has 1 fewer black node.
 - Correct by giving x an “extra black.”
 - Add 1 to count of black nodes on paths containing x .
 - Now property 5 is OK, but property 1 is not.
 - x is either *doubly black* (if $color[x] = \text{BLACK}$) or *red & black* (if $color[x] = \text{RED}$).
 - The attribute *color[x]* is still either RED or BLACK. No new values for *color* attribute.
 - In other words, the extra blackness on a node is by virtue of x pointing to the node.

Remove the violations by calling RB-DELETE-FIXUP:



Deletion

RB-DELETE-FIXUP(T, x)

while $x \neq \text{root}[T]$ and $\text{color}[x] = \text{BLACK}$

do if $x = \text{left}[p[x]]$

then $w \leftarrow \text{right}[p[x]]$

if $\text{color}[w] = \text{RED}$

then $\text{color}[w] \leftarrow \text{BLACK}$

▷ Case 1

$\text{color}[p[x]] \leftarrow \text{RED}$

▷ Case 1

 LEFT-ROTATE($T, p[x]$)

▷ Case 1

$w \leftarrow \text{right}[p[x]]$

▷ Case 1

if $\text{color}[\text{left}[w]] = \text{BLACK}$ and $\text{color}[\text{right}[w]] = \text{BLACK}$

then $\text{color}[w] \leftarrow \text{RED}$

▷ Case 2

$x \leftarrow p[x]$

▷ Case 2

else if $\text{color}[\text{right}[w]] = \text{BLACK}$

then $\text{color}[\text{left}[w]] \leftarrow \text{BLACK}$

▷ Case 3

$\text{color}[w] \leftarrow \text{RED}$

▷ Case 3

 RIGHT-ROTATE(T, w)

▷ Case 3

$w \leftarrow \text{right}[p[x]]$

▷ Case 3

$\text{color}[w] \leftarrow \text{color}[p[x]]$

▷ Case 4

$\text{color}[p[x]] \leftarrow \text{BLACK}$

▷ Case 4

$\text{color}[\text{right}[w]] \leftarrow \text{BLACK}$

▷ Case 4

 LEFT-ROTATE($T, p[x]$)

▷ Case 4

$x \leftarrow \text{root}[T]$

▷ Case 4

else (same as **then** clause with “right” and “left” exchanged)

$\text{color}[x] \leftarrow \text{BLACK}$



Deletion

Idea: Move the **extra black up** the tree until

- x points to a red & black node \Rightarrow turn it into a black node,
- x points to the root \Rightarrow just remove the extra black, or
- we can do certain rotations and recolorings and finish.


Within the **while** loop:

- x always points to a nonroot doubly black node.
- w is x 's sibling.
- w cannot be $nil[T]$, since that would violate property 5 at $p[x]$.


There are 8 cases, 4 of which are symmetric to the other 4. As with insertion, the cases are not mutually exclusive. We'll look at cases in which x is a left child.

Deletion

Case 1: w is red




- w must have black children.
- Make w black and $p[x]$ red.
- Then left rotate on $p[x]$.
- New sibling of x was a child of w before rotation \Rightarrow must be black.
- Go immediately to case 2, 3, or 4.



Deletion

Case 2: w is black and both of w 's children are black




[Node with gray outline is of unknown color, denoted by c .]

- Take 1 black off x (\Rightarrow singly black) and off w (\Rightarrow red).
- Move that black to $p[x]$.
- Do the next iteration with $p[x]$ as the new x .
- If entered this case from case 1, then $p[x]$ was red \Rightarrow new x is red & black \Rightarrow color attribute of new x is RED \Rightarrow loop terminates. Then new x is made black in the last line.



Deletion


Case 3: w is black, w 's left child is red, and w 's right child is black



- Make w red and w 's left child black.
- Then right rotate on w .
- New sibling w of x is black with a red right child \Rightarrow case 4.


Deletion

Case 4: w is black, w 's left child is black, and w 's right child is red



[Now there are two nodes of unknown colors, denoted by c and c' .]

- Make w be $p[x]$'s color (c).
- Make $p[x]$ black and w 's right child black.
- Then left rotate on $p[x]$.
- Remove the extra black on x ($\Rightarrow x$ is now singly black) without violating any red-black properties.
- All done. Setting x to root causes the loop to terminate.



Deletion

Analysis

$O(\lg n)$ time to get through RB-DELETE up to the call of RB-DELETE-FIXUP.

Within RB-DELETE-FIXUP:

- Case 2 is the only case in which more iterations occur.
- x moves up 1 level.
- Hence, $O(\lg n)$ iterations.
- Each of cases 1, 3, and 4 has 1 rotation $\Rightarrow \leq 3$ rotations in all.
- Hence, $O(\lg n)$ time.